

AI-Native Content Management for Healthcare: From Policy Text to Auditable Rules

Gaurvendra Pratap Singh Pundhir | University of Arizona

Strategic thesis: Cotiviti should treat healthcare policy content not as static documents, but as a governed operating layer that can be summarized, compared, converted into structured rules, tested, reviewed, and audited before it affects payment integrity, prior authorization, risk adjustment, or quality workflows.

Topic definition

Healthcare content management is the discipline of maintaining the clinical, financial, contractual, and operational knowledge that drives payer and provider workflows: billing and coding policies, clinical practice guidelines, payer-provider contracts, medical necessity criteria, prior authorization requirements, and payment integrity edits. In many organizations, this knowledge still lives as PDFs, manuals, policy portals, spreadsheets, and locally interpreted rules. The business challenge is not only storing the content; it is translating content changes into safe, explainable, version-controlled operations.

Why the topic matters now

Three forces make this a high-value investment area. First, the CMS Interoperability and Prior Authorization Final Rule increases pressure on impacted payers to improve prior authorization processes, provide more specific denial reasons, and implement prior authorization APIs on the required implementation timeline (Centers for Medicare & Medicaid Services [CMS], 2024, 2026). Second, HL7 Da Vinci implementation guides such as Coverage Requirements Discovery, Documentation Templates and Rules, and Prior Authorization Support are pushing coverage and documentation requirements closer to machine-readable payer-provider workflows (HL7 International, n.d.-a, n.d.-b, n.d.-c). Third, generative AI can summarize, classify, compare, and draft structured logic from long policy text, but healthcare requires governance because hallucinated or poorly traced rules can create payment errors, provider abrasion, patient delays, and compliance exposure.

Cotiviti relevance

Cotiviti's stated healthcare focus includes payment integrity, risk adjustment, quality improvement, and consumer engagement; its 2025 acquisition of Edifecs added a major interoperability capability to that context (Cotiviti, 2025a, 2025b). This creates a natural strategic opening: use AI-assisted content intelligence to connect policy knowledge with claims accuracy, risk and quality operations, and prior authorization workflows. The strongest product position is not an autonomous denial engine. It is a human-in-the-loop workbench that accelerates policy interpretation while preserving source citations, reviewer approvals, and audit trails.

Trends and opportunities

The most important trend is the shift from document management to computable policy management. A practical system could ingest a policy update, summarize the change, highlight affected criteria, draft structured rules, create regression test cases, and show the source paragraph behind each operational output. For Cotiviti, this creates four opportunities: (1) policy-change intelligence for CMS, payer, contract, and coding updates; (2) policy-to-rule conversion for payment integrity edits and prior authorization criteria; (3) auditable explanation layers for providers, compliance teams, and appeal workflows; and (4) interoperability-ready content services aligned to FHIR and Da Vinci standards.

Threats and controls

The threats are real. AI may hallucinate criteria, miss exceptions, misread conflicting policies, or produce rules that are not clinically, contractually, or regulatorily valid. Healthcare AI also raises privacy, security, fairness, and transparency issues; ONC's HTI-1 rule and the NIST AI Risk Management Framework both reinforce the need for transparency, risk management, accountability, and trustworthy design (Office of the National Coordinator for Health Information Technology [ONC], 2024; National Institute of Standards and Technology [NIST], 2023). The solution should therefore separate AI-drafted content from approved operational logic, require human review for material rules, avoid PHI unless strict controls are in place, log every transformation, and test outputs against expert-approved examples before deployment.

Strategic recommendation

Cotiviti should explore a focused MVP called PolicyLens Workbench: a policy operations intelligence layer that compares old and new healthcare policies, maps operational impact, extracts candidate rules into structured form, generates QA tests, runs synthetic claim or authorization scenarios, manages reviewer approval, and exports an audit package linking source text to output. Success metrics should include policy review cycle time, extracted-rule precision and recall against subject-matter-expert gold standards, percentage of rules with source traceability, reviewer override rate, downstream appeal or correction rate, QA coverage, and implementation readiness for FHIR-aligned workflows.

Proof-of-concept produced

PolicyLens Workbench is a React/TypeScript prototype built with synthetic data only. The app includes seven policy domains, 33 source-cited candidate rules, 23 synthetic cases, 27 changed clauses, editable scenario inputs, reviewer status controls, QA test generation, and exportable audit packages. The POC intentionally does not process PHI, does not make real medical/payment/coverage decisions, and is designed to prove the governed workflow: policy update -> version diff -> impact map -> candidate rules -> reviewer actions -> scenario testing -> QA tests -> audit-ready evidence.

Operating model

The recommended operating model is deliberately governed. AI or automated extraction should create candidate rules only; subject-matter experts approve the rule, compliance validates traceability, engineering receives structured implementation notes, and QA receives generated positive, negative, boundary, and exception tests. This separation matters because healthcare policy logic is not merely technical logic. It encodes contract terms, regulatory obligations, clinical documentation expectations, and provider-facing explanations.

Measurement plan

A pilot should measure whether the workbench reduces policy review cycle time and improves implementation quality without increasing risk. The most useful metrics are time from policy update to approved rule package, percent of rules with source traceability, QA tests generated per rule, reviewer override rate, unresolved high-risk changes, audit package completeness, and downstream correction or appeal trends. The strategic position is simple: governed AI should make healthcare policy faster to operationalize, safer to audit, and easier to explain.

References

- American Medical Association. (2025). 2025 AMA prior authorization physician survey. <https://www.ama-assn.org/system/files/prior-authorization-survey.pdf>
- Centers for Medicare & Medicaid Services. (2024, January 17). CMS interoperability and prior authorization final rule CMS-0057-F: Fact sheet. <https://www.cms.gov/newsroom/fact-sheets/cms-interoperability-and-prior-authorization-final-rule-cms-0057-f>
- Centers for Medicare & Medicaid Services. (2026, June 2). CMS interoperability and prior authorization final rule CMS-0057-F. <https://www.cms.gov/initiatives/burden-reduction/overview/interoperability/policies-regulations/cms-interoperability-prior-authorization-final-rule-cms-0057-f>
- Cotiviti. (2025a, March 31). Cotiviti completes acquisition of Edifecs. <https://www.cotiviti.com/press-release/cotiviti-completes-acquisition-of-edifecs>
- Cotiviti. (2025b). Healthcare analytics careers: Solutions overview. <https://www.cotiviti.com/careers>
- Cotiviti. (n.d.). Healthcare claim payment accuracy solutions. <https://www.cotiviti.com/solutions/payment-accuracy>
- HL7 International. (n.d.-a). Da Vinci Coverage Requirements Discovery implementation guide. <https://hl7.org/fhir/us/davinci-crd/>
- HL7 International. (n.d.-b). Da Vinci Documentation Templates and Rules implementation guide. <https://hl7.org/fhir/us/davinci-dtr/>
- HL7 International. (n.d.-c). Da Vinci Prior Authorization Support implementation guide. <https://hl7.org/fhir/us/davinci-pas/>
- National Institute of Standards and Technology. (2023). Artificial intelligence risk management framework (AI RMF 1.0). U.S. Department of Commerce. <https://nvlpubs.nist.gov/nistpubs/ai/nist.ai.100-1.pdf>
- Office of the National Coordinator for Health Information Technology. (2024). HTI-1 final rule: Certification program updates, algorithm transparency, and information sharing. <https://healthit.gov/regulations/hti-rules/hti-1-final-rule/>
- Note.** The prototype uses synthetic policy text and synthetic case inputs only; it does not process PHI and is not a clinical, payment, or coverage decision system.